

## Performance Testing

*Team Members: Thrivedh Reddy (Lead), Bhagesh T., Harshavardhan R., Vijay R., Sampath K.*

### Performance Benchmarks

Initial load: Slow (downloads DistilBERT & GPT-2 models).

Inference processing: Under 1.5s on standard CPU.

Memory usage: ~1GB RAM overhead for model weights.